



Regulatory Graph v0.6/v0.7 – Alignment Review and Recommendations

1. Graph Schema & Timeline Engine – Alignment with Mission

Domain & Jurisdictional Expressiveness: The current graph schema is designed to be domain-independent and multi-jurisdictional, aligning well with the mission of a **shared regulatory rules graph** usable across industries. The schema explicitly supports multiple domains (tax, welfare, pensions, etc.) and jurisdictions without a single “primary” country ¹. For example, node types like `:Statute`, `:Benefit`, `:Relief`, and `:Guidance` can represent rules across different domains, all linked to `:Jurisdiction` nodes (e.g. Ireland, UK, EU) ² ³. Cross-border regimes and treaties are first-class: the schema includes nodes for `:Region`, `:Agreement` / `:Treaty`, and `:Regime` to model special jurisdictions (e.g. Isle of Man, CTA) and coordination agreements ⁴ ⁵. Edges like `COORDINATED_WITH` between regimes and `TREATY_LINKED_TO` between rules and treaties connect cross-jurisdictional interactions ⁶. This expressiveness addresses the original goal of handling cross-border overrides and coordination (e.g. EU rules overriding domestic laws) ⁷.

Time-Based Reasoning: The **Timeline Engine v0.2** is tightly integrated into the schema. Time-bound constraints like lookbacks, lock-ins, deadlines, and effective periods are represented via `:Timeline` nodes and specific edges (e.g. `LOOKBACK_WINDOW`, `EFFECTIVE_WINDOW`, `LOCKS_IN_FOR_PERIOD`) ⁸ ⁹. Each rule node can attach to timeline nodes expressing its temporal constraints, which the Timeline Engine can evaluate. For instance, sections or benefits connect to `:Timeline` nodes via edges like `[:EFFECTIVE_WINDOW]` or `[:USAGE_FREQUENCY]` to denote effective dates and usage limits ¹⁰. This explicit temporal modeling satisfies the need to handle time-based questions (“What if I delay until next year?”) directly in the graph structure ¹¹. The **GraphChangeDetector** further complements this by watching for updates in Memgraph and streaming changes (node additions/updates with timestamps) to clients ¹² ¹³. This ensures the knowledge graph can be incrementally updated and queries can reflect the latest rule changes over time.

Privacy & Data Boundaries: The architecture strictly enforces that **no user-specific data enters the graph**, aligning with the privacy-first mission. The Memgraph-based rules graph is *tenant-agnostic and PII-free* ¹⁴ ¹⁵. All personal profiles, scenarios, or user inputs remain outside in the application DB (Supabase) and are only applied at query time via profile tags or scenario inputs ¹⁶. The **Graph Ingress Guard** (applied to every Memgraph write) guarantees this separation by rejecting any writes containing emails, IDs, or free-text scenarios ¹⁷. It uses schema-based whitelists and regex checks to ensure only allowed public rule content is stored ¹⁸. In effect, the **shared rules graph** contains only public law and policy knowledge, never user data or case-specific facts ¹⁵. This design upholds SOC2/GDPR principles and the original non-goal of avoiding any PII in the knowledge graph ¹⁴.

Agent-Friendliness (GraphRAG & Explanations): The schema is built to support **Graph-Retrieval-Augmented Generation (GraphRAG)** and AI agent reasoning. It emphasizes structured relationships that agents can traverse and explain. For example, relationships like `REQUIRES`, `EXCLUDES`, and `COORDINATED_WITH` explicitly link rules to conditions, mutually exclusive benefits, or cross-border equivalents ¹⁹ ²⁰. This allows an agent to fetch a focused subgraph relevant to a query (e.g. all benefits that conflict with a given relief, including timeline nodes) and present it. The agent can then

use these graph links to produce evidence-backed answers (pointing to specific sections or nodes). Indeed, every answer includes a list of `referencedNodes` from the graph for transparency ²¹. The Timeline Engine's outputs are also structured so agents can narrate temporal logic ("Benefit X has a 12-month lookback window...") based on graph data rather than hard-coded rules. Overall, the **graph schema v0.4/v0.6** provides a rich yet controlled semantic network that an LLM-based agent can safely navigate. It was a design goal to **explain interactions** (not just retrieve documents) using the graph ¹¹, which the current schema achieves by modeling rule interactions and time explicitly. Agents can retrieve a small subgraph of interest and either feed it into prompts or use it to ground their answers, fulfilling the GraphRAG vision of evidence-linked responses instead of opaque advice ¹⁹.

Extensibility & "Living Graph" Behavior: The rules graph is intended to evolve continuously, which is aligned with the mission of handling changing regulations. The schema permits adding new node types or edges without breaking existing logic, and it supports **incremental upserts** for newly discovered rules or cases ²². For instance, the presence of `:Update` / `:ChangeEvent` nodes (with edges like `AFFECTS` or `AMENDED_BY`) means changes (new laws, amendments, court decisions) can be appended over time and linked to the rules they modify ²³. Every node and edge includes `created_at/updated_at` timestamps to support change tracking and real-time patch streaming ¹². The **graph schema v0.6** added SKOS-like `:Concept` and `:Label` nodes for concept tagging, and extended properties (e.g. `last_verified_at`) to manage updates ²⁴. These choices reflect an architecture that remains *evolvable* as new domains or features (like numeric tax rate nodes, obligations, forms, etc.) are introduced ²⁵ ²⁶. In summary, the graph and timeline engine as of v0.6 appear well-aligned with the original vision: they capture complex multi-domain regulatory knowledge in a time-aware, privacy-safe graph structure that is **expressive, extensible**, and readily usable by AI agents for explanation and reasoning ⁹ ²⁷.

2. Auto-Population Mechanisms - Implementation & Alignment

Public-Only Data Ingestion: The Copilot's auto-population pipeline is carefully designed to ingest *only public regulatory data*, preserving privacy. All content added to the Memgraph graph comes from **authoritative public sources** (statutes, government websites, EU regulations, case law, etc.), and the system architecture strictly mediates these ingestion flows. For example, an external **MCP (Memgraph Control Plane) gateway** and sandboxed tools are used to fetch documents (e.g. scraping Revenue guidance or EU directives), but these tools have read-only access and cannot write directly to the graph ²⁸ ²⁹. Any new rule nodes derived from such data must go through the controlled backend service. In practice, the **GraphWriteService** is the single funnel for all graph modifications ³⁰. It wraps each write in the **Graph Ingress Guard** chain, which enforces the schema and privacy rules (rejecting disallowed fields or patterns) before committing ³¹ ¹⁷. This means even if an LLM or ingestion script tries to introduce something unexpected, it will be vetted and sanitized (e.g. stripped of any accidental personal info) prior to insertion. The guardrails (ingress aspects) are non-negotiable and auditable – they log or flag any unusual data. This architecture guarantees that auto-ingestion yields the same kind of **PII-free, public-only** knowledge that manual seeding would, thereby respecting the privacy boundaries at all times ³² ³³.

Incremental Graph Enrichment via Chat & Discovery: The system embraces a **self-populating** approach: every user conversation can trigger graph growth. In the main chat, when users ask about a concept that is not yet fully represented (say "*VAT in Ireland*" or a specific benefit name), the Copilot's agents identify those **regulatory concepts** and use LLM tools to capture them ³⁴. Detected concepts are resolved to canonical `:Concept` nodes in the graph (or created if new) through the GraphWriteService ³⁵. Synonymous terms get attached as `:Label` nodes (`HAS_ALT_LABEL` relationships), and the concept is linked to any known rules (`ALIGNS_WITH` edges to statutes, sections,

benefits, etc.)³⁵. Importantly, the system can then **enqueue background ingestion jobs** based on these concepts³⁶. For example, if “VAT” is mentioned, an ingestion agent might be scheduled to fetch the official VAT Act or Revenue guidance and then upsert relevant sections and timelines via the GraphWriteService³⁷. This mechanism allows the graph to **grow organically** from real usage, focusing on areas of interest to users. It’s an incremental, feedback-driven enrichment: new public data is added gradually, seeded by chat queries and curated by the ingestion pipeline. Crucially, each upsert remains controlled – done via aspect-guarded writes (with full validation and logging) rather than ad-hoc graph mutations. The current implementation (v0.6) has concept capture fully working – conversations emit SKOS-like metadata which is parsed and merged into the shared graph³⁸. This shows high fidelity to the principle of *learning from use*: the more the Copilot is used, the richer the shared knowledge becomes, all without ever leaking user specifics into the graph³⁹.

Controlled, Auditable Upserts (GraphWriteService + Aspects): The use of a **GraphWriteService** with an aspect pipeline for all writes provides strong guarantees of fidelity and auditability. Every insertion or update is executed via high-level methods (e.g. `upsertRule`, `upsertRegime`) that internally create a structured `GraphWriteContext` object³⁰⁴⁰. This context passes through a series of **ingress aspects** – modular functions that can validate or augment the write before it reaches the database. Baseline aspects enforce the critical rules: the schema is respected (no unknown node/edge types), only whitelisted properties are allowed, and static checks ensure no PII or tenant data is present¹⁸¹⁷. These guarantees cannot be bypassed or turned off. On top of them, there is room for **custom aspects**, which can implement additional policies or transformations (configured per deployment)⁴¹. Because this chain is executed for every write – whether triggered by an MCP ingestion bot or by an LLM extraction – the system can maintain an **audit trail** of changes. Each `GraphWriteContext` could carry metadata about its source (ingestion job ID, source URL, or conversation trigger) and even be logged by an `AuditTaggingAspect`⁴¹. This means regulators or developers can later review *who/what introduced a given node* and on what basis³³. The design even anticipates using **AI policy agents as aspects** (for example, an aspect that uses an LLM to double-check the suggested insertion for consistency) without ever relinquishing the final control that the baseline rules impose⁴²⁴³. In sum, the auto-population implementation is built with **safety and accountability** in mind: it faithfully ingests only allowable public data, enhances the graph in a stepwise manner driven by real information needs, and does so through a single, guarded entrypoint so that every modification can be verified and trusted.

3. Recommended Enhancements and Future Improvements

Building on versions 0.6 and 0.7, several enhancements can increase the system’s power and fidelity. Below we identify **high-value** improvements (incremental gains with relatively straightforward implementation) and **very high-value** improvements (more ambitious features addressing core mission goals) across three areas: graph schema, ingestion pipeline, and reasoning workflows.

3.1 Graph Schema Enhancements

High-Value Enhancements:

- **Introduce Explicit “Override/Supersedes” Relationships:** Add first-class support for modeling legal overrides, where a rule in one jurisdiction supersedes or nullifies a rule in another. While cross-border links exist (e.g. treaties coordination), an explicit `OVERRIDES` or `SUPERSEDES` edge would clarify cases like EU directives overriding domestic law. This would enable agents to immediately identify when a domestic provision is trumped by a higher authority. For example, if an EU Regulation has direct effect, linking it via `:OVERRIDES->(:Section)` of national law

gives a clear graph signal for reasoning. This addresses the cross-jurisdictional override scenarios noted in the original vision ⁴⁴ and ensures the graph can represent *hierarchy of authority*. It's a targeted schema change that would greatly help explain "Which rule takes precedence?" in multi-jurisdiction contexts.

- **Versioning and Change History on Nodes:** Enrich the schema to better handle evolving rules over time. Currently, change events and timeline edges mark when sections are amended or repealed, but an enhanced versioning system could treat major updates as distinct node versions linked together. For instance, a `:Section` node could have a property or sub-node for each major amendment, or old versions could be maintained as separate nodes with `AMENDED_BY` links ²³. Introducing a clear versioning convention (like sequence numbers or effective date ranges per node version) would make historical queries ("What did the law say in 2020 versus 2024?") easier. This is in line with schema goal #5 (support change tracking) ⁴⁵. High-value here is maintaining backward references for research – enabling the copilot to answer questions about past regulations or to show *before vs after* comparisons when rules change.
- **Strengthen Schema Constraints & Validation:** Extend the Graph Ingress Guard's schema checks to cover more nuanced constraints. For example, ensure certain relationships only occur between specific node types or under specific properties (perhaps using Memgraph triggers or by expanding the `SchemaValidationAspect`). This could include constraints like "an `:EUDirective` must link to at least one `:Section` via an implementation edge" or "`:Timeline` nodes of kind `DEADLINE` must be attached to an `:Obligation`". These business rules reduce inconsistent data entry. They can be implemented as additional whitelist rules or as a validation aspect in `GraphWriteService`. Such upfront enforcement yields a cleaner, more reliable graph for agents to use, at relatively low implementation cost (since the aspect system is already in place to add these rules).
- **GraphRAG Metadata Nodes:** To further improve agent explanations, consider adding optional helper nodes that capture *explanation structures* without altering core semantics. For example, nodes representing clusters of related rules or common justification notes could be linked via a special edge (non-public, internal use). The roadmap already hints at algorithm-derived metadata like communities or centrality ⁴⁶. Implementing Leiden community detection to group highly connected rules, for instance, and adding a `:Community` node tag with `BELONGS_TO_COMMUNITY` edges would be a high-value addition. Agents could use these groupings to say "These three provisions are commonly related" or to prioritize which parts of the graph to retrieve first (improving GraphRAG efficiency). Since these are strictly supplemental and can be toggled on/off ⁴⁷, they wouldn't compromise the core schema and can be added as needed.

Very High-Value Enhancements:

- **Multi-Domain Ontology Expansion:** Evolve the schema into a more comprehensive **ontology** that could extend beyond the initial domains and easily plug in new industries. For example, if moving into environmental regulations or healthcare compliance, the graph might need new node types (`:Permit`, `:Standard`, `:Inspection`) or new relationship types (like `COMPLIES_WITH`, `VIOLATION_OF`). A very high-value initiative is to define a clear extension mechanism for node types – perhaps a versioned schema plugin system or a taxonomy mapping. This would ensure the "domain-independent" promise truly scales to *any industry*. It could involve aligning the internal schema with widely-used standards (like SKOS for concepts, or

LegalRuleML/OWL for legal ontologies) to maximize reuse. While complex, this positions the copilot as a universal compliance graph platform.

- **Integrated Numerical Modeling (Rates & Thresholds):** Leverage the existing `:Rate`, `:Threshold`, and `:Obligation` nodes ²⁵ ²⁶ to implement actual computational logic in the graph. A high-impact improvement is to assign formulas or values to these nodes and constraints so that the graph not only knows relationships, but can be used to compute outcomes (e.g. calculate a tax based on a rate threshold). This may involve storing formula metadata or references to code that can compute using graph data. It blurs into the reasoning layer, but having the schema explicitly support calculations – for instance, linking `:Rate` nodes to specific conditions or outcomes – would enrich the question-answering. Agents could then answer “How much tax would I owe?” by traversing to the correct rate and performing a calculation rather than deferring entirely to external tools. Memgraph’s GraphQL or custom procedures could be used to execute these calculations. Embedding such capabilities is very high-value as it transforms the graph from static knowledge into an *interactive rules calculator*, all while keeping the logic transparent and explainable via the graph structure.
- **Enhanced Cross-Regime Alignment via Concepts:** Expand the use of `:Concept` nodes to actively link equivalent or similar rules across jurisdictions. Currently, concepts anchor common ideas (like “VAT”) and align with rules via `ALIGNS_WITH` ²⁴. A next step is to automatically discover and maintain **equivalence relationships**: e.g., if Ireland has a “Home Carer Tax Credit” and the UK has a roughly analogous “Marriage Allowance,” linking those via a shared concept or a new `EQUIVALENT_TO` edge can let the system transfer knowledge across regimes. This might involve an offline alignment process or an AI-assisted matching of graph substructures. The value here is extremely high for cross-border use cases – it would allow a question about one country’s rule to prompt the agent to mention the counterpart in another country. Essentially, the graph becomes not just a silo per jurisdiction but a connected lattice of regimes, fulfilling the “*explain equivalent rules across jurisdictions*” vision ⁴⁸. Implementing this might require new schema elements (like a property indicating a concept’s domain scope or confidence in equivalence) and likely an agent/pipeline to find the alignments. But once in place, it significantly enhances the copilot’s ability to handle cross-regime questions and multi-jurisdiction scenarios gracefully.

3.2 Auto-Ingestion Pipeline Enhancements

High-Value Enhancements:

- **LLM-Based Quality Assurance for Ingestion:** Introduce a second-pass validation using LLMs or rule-based checks on data just ingested. For instance, after an MCP tool parses a regulatory document into graph nodes, an LLM agent could be asked to verify the consistency of those nodes with the source text (in a read-only manner). Any discrepancies or low-confidence extractions could be flagged for review or logged. This would act as a form of *runtime supervision*, catching errors in parsing (e.g. mis-identified section headings or mis-classified provisions) before they propagate. Since all ingestion goes through the GraphWriteService, this QA could be implemented as a custom Graph Ingress Aspect that pauses writes pending an asynchronous verification. This improvement is relatively contained (does not require changes to user experience) but boosts the fidelity of the auto-populated graph.
- **Policy Agent Ingress Aspect:** Leverage the architecture’s support for AI policy agents in the write pipeline ⁴² by implementing a concrete **policy-aspect agent**. For example, an aspect could use an AI model to decide if a proposed insert aligns with organisational policy (e.g. only

ingest certain jurisdictions or ignore very old laws). It could also add metadata tags (like `source_trust_level` or `extraction_confidence`) into the write context's metadata for audit purposes ⁴¹. A simple initial version might be a check that the source of the data is on an approved list (to ensure fidelity to **public-only** data). Over time, more sophisticated agents could detect subtle PII (beyond regex, using context) or enforce jurisdiction-specific data residency rules. Implementing such agents as aspects is high-value because it enhances trust and compliance of the ingestion process itself, and the framework is already in place to add these without weakening baseline guarantees ⁴³.

- **Incremental Ingestion & Refresh Jobs:** Build on the “living graph” concept by scheduling **regular refresh and incremental ingestion tasks**. For example, a high-value enhancement is an ingestion crawler that periodically checks official sources for updates (new legislation, changes in tax rates) and automatically updates the graph. This might involve storing last fetched timestamps in the `metadata` of nodes (the schema already allows a `last_verified_at` on concept nodes ²⁴, and similar could be done for others). By having a daemon process or cron job for each integrated data source, the copilot remains up-to-date without waiting for a user to ask about a topic. The architecture’s change detection (`GraphChangeDetector`) could then broadcast these updates to active sessions ⁴⁹. This proactive ingestion pipeline turns the graph into a continuously curating knowledge base. It’s not as complex as full real-time web crawling (since targets can be well-defined, like an annual Finance Act publication), but yields substantial value by keeping the knowledge current and reducing stale information.

Very High-Value Enhancements:

- **Regulatory Document Ingestion via Fine-Tuned Models:** Invest in fine-tuning or training domain-specific NLP models to parse regulatory texts into graph structures more accurately. Currently, LLMs (likely general ones via the `LlmRouter`) are used to assist parsing and concept extraction. A very high-value step would be developing a specialized model (or prompt chain) that understands legal document structure – sections, subsections, definitions – and populates the graph accordingly with minimal human post-processing. For example, a fine-tuned model on tax law documents could directly output JSON for nodes/edges to be fed into `GraphWriteService`. This could dramatically speed up onboarding a new corpus of regulations. It’s a significant effort (data gathering and training required), but it pays off by enabling faster expansion into new domains or jurisdictions with confidence in the structure of ingested data. Such a model could be integrated as an MCP tool or within the sandbox environment to keep it isolated and safe, similar to how code execution tools were added in v0.7 ⁵⁰.
- **Ingestion Simulation & Dry-Run Mode:** Develop a “dry-run” or simulation mode for ingestion jobs where the system processes a document and produces a diff or report **without immediately modifying the live graph**. This would be akin to a CI pipeline for knowledge ingestion. A policy team or developer could review the proposed additions/changes (perhaps via a temporary visualization or summary of new nodes/relationships) and approve or reject them. This enhancement is very high-value for governance and auditability, especially as the graph scales to many domains. It introduces a human-in-the-loop checkpoint for quality. Implementing this could involve a staging Memgraph or simply a transaction that is rolled back after diff generation. The Graph Ingress Guard aspects could be reused in this mode to show what *would* be blocked or altered. While it introduces some overhead, it would ensure that large-scale ingestions (say ingesting an entire new Act or thousands of case laws) can be sanity-checked – aligning with the need for an auditable, **controlled growth** of the graph.

- **Feedback Loop from Usage (Learning which gaps to fill):** Enhance the ingestion pipeline with a mechanism that analyzes user queries and agent performance to identify high-value content gaps. For instance, if multiple users ask about a concept that the copilot struggles with (perhaps the answers frequently fall back to the fallback agent or have low confidence), the system should flag this concept or area for targeted ingestion. This could manifest as an admin dashboard of “requested concepts not well covered” or automatically spawning an ingestion job for a specific term. By quantifying concept frequency and success rates (using conversation logs and the `referencedNodes` data), the pipeline can **prioritize what to ingest next**. This turns user interactions into an active learning signal. It’s very high-value because it ensures the graph’s growth is driven by actual demand, thereby improving the copilot’s utility in the areas users care about most. Implementing this ties together the conversation context store and ingestion scheduler – effectively a recommendation system for the graph. Over time, this could even be extended to allow power-users to confirm suggestions (“Yes, please fetch more on X”), making the copilot more participatory.

3.3 Reasoning Workflow Enhancements

High-Value Enhancements:

- **Timeline Comparison Tools:** Leverage the timeline engine and the emerging Scenario Engine to allow side-by-side comparisons of scenarios or dates. A concrete high-value feature is a “**timeline diff**” view or query, where a user (or agent) can ask “What’s the difference in outcome if I do X now versus in 5 years?” The system can then evaluate two scenarios – one with the event now, one with it later – and highlight differences in eligibility or obligations. Under the hood, this uses the timeline engine to compute rule applicability at different times ⁵¹ ⁵². Surfacing it might mean an agent that takes two snapshot dates as input and queries the graph for each, then returns a structured comparison. Even without a full UI, just enabling an **agentic answer that enumerates changed entitlements or triggered rules between two timelines** adds a lot of value. It directly addresses users’ mental “what-if” analysis in temporal terms, using the data already in the graph.
- **Cross-Regime “Diff” Explanations:** Similar to timeline diffs, provide an agent capability to compare regulatory regimes or jurisdictions. For example, “How would this situation differ if I were in Northern Ireland instead of Ireland?” The groundwork exists in the graph (with profile tags and jurisdiction nodes), but a specialized reasoning routine can make it explicit. A high-value enhancement is an agent that, given two jurisdiction inputs, pulls the relevant rules for each and then identifies which benefits/reliefs are available in one but not the other, which obligations change, etc. This could utilize the shared concept linking or treaty links to find corresponding rules. Essentially, it automates cross-jurisdictional alignment analysis – something currently done manually by experts. Implemented as a tool, it could output a table of differences that the LLM can then narrate. This feature would capitalize on the graph’s multi-jurisdiction modeling to fulfill the promise of explaining *interactions across IE/UK/EU regimes* ⁵³. Even as a simple list of comparative points, it would greatly help users with cross-border questions.
- **Jurisdictional Override Alerts:** Augment the reasoning workflow so that whenever an answer involves a national rule that is linked to a higher-order rule (EU law or treaty), the agent automatically includes a note or explanation about that relationship. This can be a subtle but high-value improvement in answer quality. For instance, if a user asks about a UK benefit that is constrained by a bilateral agreement, the answer should mention the agreement’s effect. Concretely, implement logic in the agent (or as a post-processing step on the answer) that checks if any referenced node has a `TREATY_LINKED_TO` or similar override edge ⁵⁴. If so,

append a clarification like “Note: This is subject to the \${agreementName} between \${countries} which may override some conditions.” By systematically including these override alerts, the copilot provides a more **complete and context-aware explanation**. This is especially helpful for edge cases (like those involving Northern Ireland and Brexit-related arrangements, etc.). It’s relatively easy to implement given the data is in the graph; it just needs to be made part of the answer assembly logic.

Very High-Value Enhancements:

- **Full Scenario Engine Implementation & UI:** Completing and extending the Scenario Engine is perhaps the most impactful next step. As noted in the project status, the scenario/what-if engine is the only major feature of v0.6 not yet implemented ⁵⁵ ⁵⁶. Finishing this (per spec v0.1) and creating a user-friendly interface for it would allow users to run **multi-step what-if simulations** with confidence. The engine would take structured scenarios (sequences of life events, profile changes over time) and evaluate applicable rules at each step ⁵¹ ⁵², returning outputs like “eligible vs ineligible benefits” and flags for lock-ins or deadlines ⁵² ⁵⁷. A UI could let users compare two scenarios (e.g. *Scenario A: stay a sole trader; Scenario B: incorporate a company in 2026*) and see differences side by side. This directly addresses complex, longitudinal questions and embodies the “research partner” vision by giving evidence-based projections. Moreover, hooking the Scenario Engine up to an **explanation agent** (as envisioned for a What-If Scenario Agent) is very high-value ⁵⁸. The agent would use the structured scenario evaluation to narrate *why* one scenario yields certain benefits or not, referencing specific rules. This synergy of computation + explanation would be a standout capability. Given it is already planned (ADR and roadmap for v0.7+), prioritizing it will yield a copilot that not only answers static questions but also guides users through prospective decision-making – a key original goal.
- **Eligibility Explorer (Deterministic Q&A):** Implement a deterministic rules evaluation mode (the “Eligibility Explorer”) for straightforward profile queries. This enhancement would allow the system to answer questions like “Am I eligible for X benefit?” through a rule-based evaluation rather than an open-ended LLM response. By codifying certain eligibility criteria into the graph (via `:Condition` nodes and edges like `REQUIRES` ⁵⁹) and having an engine walk through those, the copilot can output a clear Yes/No plus reasoning. The roadmap highlights this as a future use case ⁶⁰. A very high-value deliverable would be a form-driven interface where a user inputs a few facts (age, income band, etc.), then the engine checks those against the graph’s conditions to produce an eligibility verdict with cited rules. This taps into the graph’s structured nature to provide *conclusive answers* where possible, enhancing trust. It also complements the conversational agent: simpler queries get deterministic answers, while complex ones still use the LLM+graph approach. Implementing the Eligibility Explorer solidifies the Copilot’s utility for common compliance questions by providing quick, auditable outcomes (and it can be internally reused by agents as a tool).
- **Advanced GraphRAG and Explanation Generation:** Finally, invest in the next generation of **Graph-based Retrieval-Augmented Generation** within the copilot. This involves both improving how relevant graph context is selected and how answers are formulated from it. A very high-value improvement would be to use graph algorithms (like centrality or community detection ⁴⁶) to smartly pick which parts of a large graph to feed into the LLM for a given query – essentially focusing the agent on the most important nodes first. Additionally, one could implement templated explanation modules for common patterns. For example, if two benefits have an `EXCLUDES` relationship, the answer could automatically include a sentence like “Note: You cannot combine Benefit A with Benefit B ²⁰.” This reduces the cognitive load on the LLM and ensures consistency in explanations. Another idea is to maintain a library of explanation

snippets for specific nodes (like a brief on what each benefit is), which the agent can pull in as needed (somewhat akin to few-shot prompts but stored in the graph as attributes). Pushing the envelope on GraphRAG will result in answers that are not only evidence-backed but also *structured and comprehensive*, covering edge conditions and interactions systematically. Given the user base (advisors, researchers) ⁶¹ ⁶², this level of depth and reliability in answers is crucial. It's a complex enhancement involving research into prompt engineering and graph algorithms, but it directly drives the core experience of the Regulatory Intelligence Copilot as a "graph-backed research partner."

References:

1. Copilot Concept v0.6 – Problem, Vision, and Goals ⁶³ ⁶⁴
2. Graph Schema Spec v0.3/v0.4 – Multi-domain, Multi-jurisdiction design ¹ ⁹
3. Graph Schema Spec v0.6 – Goals (cross-border, change tracking, PII-free) ²⁷ ⁶⁵
4. Graph Schema Spec – Privacy & PII constraints summary ⁶⁶
5. Graph Schema Spec – Agent-friendly design and relationships ¹⁹ ²⁰
6. Graph Schema Spec – Update/ChangeEvent nodes and timeline edges ²³ ¹⁰
7. Graph Ingress Guard Spec – All writes via service, aspect validation ³¹ ³⁰
8. Graph Ingress Guard Spec – Baseline PII and schema enforcement ¹⁸ ¹⁷
9. Architecture v0.6 – Use of MCP ingestion jobs and GraphWriteService ³⁷ ⁶⁷
10. README – Self-populating graph via concept capture from chat ³⁶
11. Copilot Concept v0.6 – Self-population vision and never storing user data ³⁴ ¹⁴
12. Copilot Concept v0.6 – Write-guarded graph, read-only to agents ⁶⁸
13. Copilot Concept v0.6 – Concept capture workflow (GraphWriteService) ³⁵
14. Roadmap v0.6 – Phase 4 (Graph ingestion, change detection goals) ⁶⁹ ⁷⁰
15. Implementation Status – v0.6 scenario engine gap (future work) ⁵⁵ ⁵⁶
16. Roadmap v0.6 – Future Use Case 7.6 (Scenario Engine & comparison) ⁷¹ ⁵⁸

¹ ² ⁸ ⁹ ¹⁰ ¹¹ ¹⁵ ¹⁶ ¹⁹ ²² ²³ ⁴⁸ ⁶⁶ **schema_v_0_4.md**
https://github.com/airnub-labs/regulatory-intelligence-copilot/blob/0234f8cca28c51f659d47e58253cdb7c61583cea/docs/architecture/graph/archive/schema_v_0_4.md

³ ⁴ ⁵ ⁶ ¹² ²⁰ ²⁴ ²⁵ ²⁶ ²⁷ ⁴⁵ ⁵⁴ ⁵⁹ ⁶⁵ **schema_v_0_6.md**
https://github.com/airnub-labs/regulatory-intelligence-copilot/blob/0234f8cca28c51f659d47e58253cdb7c61583cea/docs/architecture/graph/schema_v_0_6.md

⁷ ¹⁴ ²¹ ³⁴ ³⁵ ³⁹ ⁴⁴ ⁶³ ⁶⁴ ⁶⁸ **concept_v_0_6.md**
https://github.com/airnub-labs/regulatory-intelligence-copilot/blob/0234f8cca28c51f659d47e58253cdb7c61583cea/docs/architecture/copilot-concept/concept_v_0_6.md

¹³ ²⁸ ³⁷ ⁶⁷ **architecture_v_0_6.md**
https://github.com/airnub-labs/regulatory-intelligence-copilot/blob/0234f8cca28c51f659d47e58253cdb7c61583cea/docs/architecture/architecture_v_0_6.md

¹⁷ ¹⁸ ³⁰ ³¹ ³² ³³ ⁴⁰ ⁴¹ ⁴² ⁴³ **graph_ingress_v_0_1.md**
https://github.com/airnub-labs/regulatory-intelligence-copilot/blob/0234f8cca28c51f659d47e58253cdb7c61583cea/docs/architecture/guards/graph_ingress_v_0_1.md

²⁹ ⁵⁰ **architecture_v_0_7.md**
https://github.com/airnub-labs/regulatory-intelligence-copilot/blob/0234f8cca28c51f659d47e58253cdb7c61583cea/docs/architecture/architecture_v_0_7.md

[36](#) [53](#) [61](#) [62](#) README.md

<https://github.com/airnub-labs/regulatory-intelligence-copilot/blob/0234f8cca28c51f659d47e58253cdb7c61583cea/README.md>

[38](#) [55](#) [56](#) V0_6_IMPLEMENTATION_STATUS.md

https://github.com/airnub-labs/regulatory-intelligence-copilot/blob/0234f8cca28c51f659d47e58253cdb7c61583cea/docs/development/V0_6_IMPLEMENTATION_STATUS.md

[46](#) [47](#) [49](#) [58](#) [60](#) [69](#) [70](#) [71](#) roadmap_v_0_6.md

https://github.com/airnub-labs/regulatory-intelligence-copilot/blob/0234f8cca28c51f659d47e58253cdb7c61583cea/docs/governance/roadmap/roadmap_v_0_6.md

[51](#) [52](#) [57](#) spec_v_0_1.md

https://github.com/airnub-labs/regulatory-intelligence-copilot/blob/0234f8cca28c51f659d47e58253cdb7c61583cea/docs/architecture/engines/scenario-engine/spec_v_0_1.md