



Regulatory Intelligence Copilot: Mission, GraphRAG Approach, and Ralph Wiggum Orchestration Potential

Project Mission and Current Architecture

The **Regulatory Intelligence Copilot** is a research-focused assistant for exploring how various rules (tax, welfare, pensions, CGT, EU regulations, etc.) interact, without giving formal advice ¹. It features a **chat-first** user interface where queries are answered by traversing a **shared rules graph** (built on Memgraph) instead of doing ad-hoc web searches ². In practice, regulations and conditions are modeled as graph nodes/edges, and the engine retrieves answers via **graph-based reasoning** over this knowledge base ³. The chat engine streams answers to the user in real time *while capturing structured concepts and evidence behind the scenes* ². This means as it answers questions, it also logs key regulatory concepts or rule references relevant to the query.

The system uses multiple specialized **expert agents** (for specific domains like Irish company tax, social welfare, investments, etc.) and a **Global Regulatory Agent** that can reason across all domains for complex cross-domain questions ⁴. Under the hood, it's a Next.js 16 + Node.js 24 application with a single `/api/chat` endpoint. User conversations and private data persist in a Postgres/Supabase database, whereas the Memgraph graph remains a **global, PII-free knowledge graph** of public rules ⁵. This design ensures privacy boundaries – user-specific context is kept out of the shared graph ⁵. Notably, v0.6 of the architecture introduced "**concept capture**", where the chat's LLM can emit metadata about relevant concepts during an answer. The backend then adds those concepts as nodes in the graph (using a SKOS-like schema) and updates the conversation context with references to those nodes ⁶. In essence, the graph can **self-populate with new concepts** mentioned in chats, creating a feedback loop where the assistant's insights enrich the knowledge graph over time ⁶. This approach – sometimes called **GraphRAG (Graph-based Retrieval Augmented Generation)** – goes beyond simple document lookup by using the structured graph to fetch precise rule snippets or relationships needed to answer questions ⁷.

"Ralph Wiggum" Orchestration Technique Overview

Ralph Wiggum (a tongue-in-cheek name from a Simpsons character) refers to an **autonomous AI coding agent loop technique** popularized by Geoffrey Huntley and others. A **Ralph Wiggum-style orchestrator** runs an AI agent continuously in a tight loop, having it tackle a complex task through many small iterative steps ⁸. Instead of a single prompt-response cycle, the agent keeps working until it either completes the task or hits certain limits. This orchestration pattern was designed for *AI-driven software development*: the agent writes code, tests it, fixes bugs, and so on in successive iterations, guided by an external loop manager. Essentially, the orchestrator repeatedly feeds the agent a prompt (including the latest project state and any instructions), the agent produces an output or code changes, and the orchestrator evaluates if the task is done or if another iteration is needed.

A reference implementation of this is **Ralph Orchestrator** by Mikey O'Brien, which provides a production-ready harness for such loops ⁸ ⁹. It supports multiple AI coding agents (Anthropic

Claude, OpenAI “Q”/GPT, Google Gemini, or any Agent Protocol-compliant model) and includes features like automatic agent selection, web search integration for the agent, error handling with retries, and state persistence across runs ⁹. One key feature is **checkpointing and an agent scratchpad**: the orchestrator saves the agent’s intermediate work and context to a scratchpad (a Markdown file, e.g. `.agent/scratchpad.md`) so that the agent remembers progress from one iteration to the next ¹⁰ ¹¹. This scratchpad acts as the agent’s “working memory,” accumulating partial code, notes, and plan updates throughout the loop. By looping with a persistent scratchpad, the agent can handle tasks that exceed a single prompt’s capacity – for example, generating and refining thousands of lines of code or performing multi-step reasoning over hours. The net result is an **autonomous coding loop** that can produce complex outputs (even entire codebases or extensive analyses) given high-level instructions, with minimal human intervention.

Potential Value of a Ralph Wiggum-Style Orchestrator for This Project

Integrating a Ralph Wiggum-style orchestrator into the Regulatory Intelligence Copilot could be a **high-value enhancement**, depending on the goals. Currently, the project’s focus is on interactive Q&A and research assistance (answering user questions with the help of the rules graph). Each user query triggers a single compliance engine run that fetches relevant graph data and composes an answer in one go. There isn’t an existing need for multi-iteration loops in the user-facing chat flow – *unless* we envision more complex autonomous tasks like scenario analyses or long-running research sessions. However, the development team *is* leveraging AI agents to build the project itself. In fact, the repository provides an official **coding agent prompt (`PROMPT.md`)** and guidelines for AI contributors ¹², implying that AI (like Claude or GPT-4) is already used to implement features within the repo. Currently those agents follow a structured prompt and the project’s extensive guidelines, but each coding session may be relatively bounded. A Ralph Wiggum orchestrator could take this a step further by allowing an AI agent to work continuously on coding tasks (or documentation, data ingestion, etc.) overnight or over long durations without human oversight – potentially accelerating development.

The value of such an orchestrator integration can be viewed in two ways: **(1) Internal development automation**, and **(2) Enhancing the product’s AI capabilities**. For internal development, using the orchestrator with a more powerful coding model (e.g. GPT-4 Code Interpreter or future Claude improvements) could dramatically speed up implementing the roadmap. The orchestrator’s loop, combined with the repo’s detailed test cases and bug-pattern guides, might let an AI fix tricky bugs or implement new features iteratively. The project has identified many “insidious bug patterns” and has a rigorous checklist for code changes ¹³ ¹⁴. An orchestrator could ensure the AI agent iteratively runs tests and checks against these patterns in each loop, only stopping when code meets all criteria. This could yield higher-quality contributions from the AI agent, reducing manual rework. In short, for **internal use**, the Ralph-style loop would harness the AI more fully – rather than one-off prompts, the AI could sustain multi-hour reasoning and coding sessions to deliver complex updates (this matches the trend noted by observers that “agents now sustain multi-hour reasoning” in autonomous loops, as per the Ralph Wiggum technique reports).

For **enhancing the product itself**, we might consider future features where the Copilot agent performs multi-step reasoning autonomously on behalf of the user. Imagine a “Regulation Simulator” or **scenario engine** that the user can task with exploring a complex what-if scenario: an orchestrator could let the agent break the scenario into sub-tasks (gather relevant rules, apply timelines, cross-verify outcomes) in a loop until it has a comprehensive answer or simulation result. This aligns with the project’s roadmap mentioning scenario and what-if capabilities as future extensions ¹⁵. A lightweight orchestrator could manage these multi-step analytical tasks behind the scenes. However, such a use in production would

need careful control (to avoid the agent looping endlessly or going off-track) and strong guardrails (privacy and safety checks on each iteration, since the Copilot must not make external calls or leak data without oversight ⁵). Given that the architecture already allows sandboxed tool execution (via E2B sandboxes and an **Egress Guard** for external calls ¹⁶), a looping orchestrator could potentially be contained within the same safety framework. In summary, adding a Ralph Wiggum-style orchestrator could be highly beneficial internally (accelerating development and ensuring alignment with design specs), and it opens the door for richer autonomous analysis features in the Copilot's future versions. The caveat is that it adds complexity – the team would need to maintain the orchestrator logic and ensure the AI doesn't produce undesirable outcomes over long loops – but the proven implementations (like O'Brien's orchestrator) and extensive test coverage ⁹ suggest these challenges can be managed.

From GraphRAG to Structured Agent Memory

A core strength of the Copilot's implementation is its **graph-backed knowledge base**. The regulatory compliance graph schema is carefully designed to model laws, benefits, conditions, timelines, etc., with explicit relationships and properties ¹⁷ ¹⁸. This makes retrieval precise and supports rich reasoning (for example, understanding that a certain tax benefit *excludes* a welfare payment due to a condition node linking them). Currently, this graph is used in a **retrieval-augmented generation** manner – i.e., when a question comes in, the agent (LLM) queries Memgraph (through the `GraphClient`) to fetch relevant nodes and their connections, and uses that information to compose an answer ¹. This is what the user described as the graph "being used as GraphRAG."

However, the team is **exploring structured agent memory**, which goes beyond using the graph just for static facts. An early example of structured memory in v0.6 is the **conversation context and concept capture** mechanism. Whenever the Copilot answers a question, it doesn't just return a textual answer – it also returns a list of `referencedNodes` (the graph node IDs that were actually used as evidence) ¹⁹. The system stores a running **ConversationContext** for each chat session (in Postgres) that includes these active node references and other state ⁶. This means the next question can be answered in context – the agent knows which specific rules or concepts have already been discussed. Additionally, through the concept capture tool, the agent can introduce new concept nodes to the graph if the user's query brings up a novel idea or a not-yet-modeled rule ⁶. In effect, the graph "learns" and expands during conversation, and the conversation state is maintained in a structured way (graph IDs rather than raw text history). This structured approach is more robust than a plain chat history: it anchors context to canonical knowledge nodes.

Given this design, it's natural to ask if the **same graph-based approach could serve as an AI agent's long-term memory** – for example, replacing the free-form `scratchpad.md` used by Ralph Wiggum loops with a more structured, queryable memory store. The idea would be to have the coding agent (in an autonomous loop) record its progress, plans, and findings in a graph data structure instead of (or in addition to) a scratchpad text. The **regulatory compliance schema** as it stands is specialized for laws and rules; it explicitly is *not* intended as a general knowledge graph ²⁰. So, in its current form, it doesn't have node types for "Coding Task" or "Function Implemented" or "Test Result". It was scoped to regulatory knowledge and deliberately excludes user-specific or incidental data ²⁰ (for privacy and focus). That said, the **principles and mechanisms** of the compliance graph could be extended to an "**agent memory graph**." In the future, one could introduce a parallel schema (or an extension of the existing one) to represent the coding agent's knowledge domain. For example, nodes could represent **tasks, code modules, bug fixes, and design decisions**, with relationships linking them (e.g., a "Task" node might connect to "CodeSnippet" nodes it created, or a "Bug" node that was resolved). The agent, during each iteration of the Ralph loop, could then update this graph: instead of appending text to a

scratchpad, it might call a tool that adds a node or link (say, “Task X – completed” or “Function Y – implemented using approach Z”).

Using a structured memory graph for the coding agent offers several potential advantages:

- **Persistent Organizational Memory:** The graph would serve as an “organizational memory” – a lasting record of what the agent has attempted or learned, which is easier to query and analyze than a monolithic markdown file. For instance, one could query “which tasks did the agent complete overnight and how are they related?” or visualize the dependency graph of tasks.
- **Context Management:** Rather than feeding the entire scratchpad to the agent every loop (which may grow large and hit context length limits), the orchestrator could fetch just the relevant subgraph as context. For example, if the agent is now working on front-end code, the orchestrator can query the memory graph for nodes related to that (recent changes in the front-end component) and supply those to the prompt. This is analogous to how the Copilot fetches only the pertinent rules from Memgraph for a given question, keeping the prompt focused.
- **Multi-Agent Collaboration:** A structured memory graph could enable multiple agents or tools to collaborate. One agent could specialize in generating code, another in testing or verifying, both reading/writing to the same memory graph of the project state. This is harder to coordinate with a plain scratchpad, but a graph could let them divvy up nodes (e.g., one agent marks a “TestCase” node as failing, another picks it up to fix the code).
- **Long-Term Scaling:** Over the long run, a graph can scale to store a huge web of information (e.g., dozens of tasks, hundreds of code snippets, their relationships and status) that an agent can navigate via queries. This could support very large coding projects or regulatory analyses that persist across sessions, whereas a scratchpad might become unwieldy.

The idea is admittedly forward-looking. In practice, implementing an **agent memory graph** would require careful design. We’d likely **not merge it directly into the regulatory rules graph**, to maintain the principle that Memgraph holds only public regulatory knowledge. Instead, the project could maintain a separate Memgraph (or a separate sub-graph namespace) for agent memory, or use another graph database optimized for transient data. The existing compliance graph schema could inform the new schema – for example, borrowing the timestamping and versioning approach (so each memory node has `created_at`, `updated_at` for time-travel debugging) and the concept of stable IDs. The team’s expertise in graph design (e.g., the emphasis on SKOS concepts for consistency ²¹ and on change tracking ²²) would translate well to designing a coding memory graph that remains coherent and queryable over time.

Importantly, the current architecture already supports **tools and extensions** that could facilitate this. The Copilot’s design allows custom tools (via the MCP tool interface) that the agent can invoke. One could create a tool like `MemoryGraphWriter` which the agent calls with a JSON payload to insert or update nodes representing its scratchpad data. Similarly, a `MemoryGraphQuery` tool could let the agent retrieve specific info (e.g., “get all incomplete tasks” or “fetch the content of node X which is the latest plan for Y”). By using such tools, an advanced coding agent (say GPT-4 or a fine-tuned domain model) could effectively maintain a live knowledge graph of its own progress. This would be a more **structured and AI-accessible memory** than a plain text file.

So, **could the regulatory compliance schema support a coding agent’s memory in the future?** In principle, **yes – the same concepts of structured, graph-based knowledge management can be applied to the coding domain**, though it would likely mean extending or parallelizing the schema to fit that domain. The Copilot’s developers have already shown interest in richer agent memory: the user indicates they are “*currently using [the graph] as GraphRAG but exploring structured agent memory.*” This suggests an active line of research in the project. Implementing it would make the Copilot not just a consumer of a static knowledge graph, but also an example of an AI system that **maintains and updates its knowledge autonomously** in a structured way. If successful, this could greatly enhance the capabilities of a “Ralph Wiggum” orchestrator powered by a more capable coding agent than Claude

- because that agent would have a memory that is both persistent and logically organized, rather than a long scroll of text.

Conclusion

In summary, the Regulatory Intelligence Copilot is an innovative platform marrying **LLM capabilities with a structured graph knowledge base** to assist with complex regulatory queries ¹. Its mission is to surface relevant rules and interactions in a user-friendly way while maintaining strict privacy and accuracy boundaries. The latest implementation (v0.6) solidifies a graph-backed architecture with concept capture and conversation context, essentially pioneering GraphRAG in a real-world domain ⁶ ⁷.

Integrating a **Ralph Wiggum-style orchestrator** could be a very fruitful enhancement. Such an orchestrator aligns with the project's exploration of agentic architectures ⁷, offering the ability to handle longer, more complex tasks through iterative AI loops. It would be particularly high-value for accelerating development via AI (letting a coding agent adhere to the project's guidelines and churn through improvements), and could open up new product features like autonomous scenario analysis. Meanwhile, the **regulatory compliance graph schema**, while domain-specific, provides a strong inspiration for how to manage structured memory. With future extensions, a graph-based memory for the coding agent could replace the simplistic scratchpad approach ¹¹, yielding better context management and persistence for AI-driven development. In effect, the project's knowledge graph might evolve into or be accompanied by an "AI memory graph," enabling a more powerful coding agent (e.g. GPT-5 or beyond) to build and reason in a way that scales with complexity.

Overall, the marriage of **GraphRAG, orchestrated agent loops, and structured memory** represents a cutting-edge path. It reinforces the Copilot's core theme: combining knowledge structure with advanced AI reasoning for higher-value outcomes. By proceeding in this direction, the team can ensure that every new capability – from automated coding to richer user assistance – stays aligned with the principle of "*everything, optimally where it's of highest value*," delivering robust features without compromising the careful design of the system. The exploratory answer, therefore, is **yes**: adopting a Ralph Wiggum orchestrator and evolving the schema for agent memory are promising steps that could significantly enhance the project's effectiveness and intelligence in the near future ²⁰ ¹¹.

Sources:

- Regulatory Intelligence Copilot README and docs ¹ ⁶ ⁷ ²⁰
- "Ralph Orchestrator" (Ralph Wiggum technique implementation) documentation ⁸ ⁹ ¹¹

¹ ² ³ ⁴ ⁵ ⁶ ⁷ ¹² ¹⁵ ¹⁶ ¹⁹ README.md

<https://github.com/airnub-labs/regulatory-intelligence-copilot/blob/e413cface5b37f30283ad2b7c806de79a8781b24/README.md>

⁸ ⁹ ¹⁰ ¹¹ GitHub - mikeyobrien/ralph-orchestrator: An improved implementation of the Ralph Wiggum technique for autonomous AI agent orchestration

<https://github.com/mikeyobrien/ralph-orchestrator>

¹³ ¹⁴ agents.md

<https://github.com/airnub-labs/regulatory-intelligence-copilot/blob/e413cface5b37f30283ad2b7c806de79a8781b24/docs/agents.md>

[17](#) [18](#) [20](#) [21](#) [22](#) **schema_v_0_6.md**

https://github.com/airnub-labs/regulatory-intelligence-copilot/blob/e413cface5b37f30283ad2b7c806de79a8781b24/docs/architecture/graph/schema_v_0_6.md